

Assignment - 3

Credit Risk Modelling – Loan Classification

Submitted By Mitesh Khadgi

Problem:

The Loan Classification project is designed to help you understand and apply machine learning technique to classify loan applicants based on their likelihood of loan repayment. This project will you hands-on experience in data pre-processing, feature engineering, model training, evaluation, and interpretation. CSV file is shared on whatsapp group.

Solution:

OUTPUT is mentioned in a separate “output.txt” file, and plots/figures are also uploaded:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import xgboost
from xgboost import XGBClassifier
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
import warnings
warnings.filterwarnings('ignore')
df = pd.read_csv('loan_detection.csv')
df.head()
df.shape
df.columns
```

```

df.info()

#Imbalance Data
df['Loan_Status_label'].value_counts()

len(df)

print(len(df[df['Loan_Status_label'] == 0])/len(df)*100)

print(len(df[df['Loan_Status_label'] == 1])/len(df)*100)

plt.pie(df['Loan_Status_label'].value_counts(), autopct = "%1.2f%%", labels = ['Loan Not Applicable',
'Loan Applicable'])

sns.countplot(x = 'Loan_Status_label', data = df)

#Missing Data
df.isnull().sum()

df.isnull().mean()*100

sns.displot(df['previous'], label = df.previous.skew(), color = 'r')

plt.legend()

sns.displot(df['no_previous_contact'], label = df.no_previous_contact.skew(), color = 'c')

plt.legend()

sns.displot(df['not_working'], label = df.not_working.skew(), color = 'g')

plt.legend()

print(df['previous'].mean())

print(df['no_previous_contact'].mean())

print(df['not_working'].mean())

df.isna().sum()

for i in df.columns:
    if df[i].isna().sum() > 0:
        df[i].fillna(df[i].mean(), inplace = True)

df.isna().sum()

#Duplicate Data
df.duplicated().sum()

df[df.duplicated()]

#Outlier Treatment
df.describe()

```

```

Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
print(Q1 - 1.5*IQR)
print('\n')
print(Q3 + 1.5*IQR)
df[~((df < (Q1 - 1.5*IQR)) | (df > (Q3 + 1.5*IQR))).any(axis = 1)]
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
data = df[~((df < (Q1 - 1.5*IQR)) | (df > (Q3 + 1.5*IQR))).any(axis = 1)]

#Feature Selection
data.corr()

#plt.figure(figsize = (12, 7))
plt.figure()
sns.heatmap(data.corr(), annot = True)

# data.corr['Loan_Status_label']

#Model Building

#Separate your Independent and Dependent data
data.head()
data.columns

# X = data[['previous', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity', 'Organic_carbon',
'not_working', 'Turbidity']]

X = df[['age', 'campaign', 'pdays', 'previous', 'no_previous_contact', 'not_working', 'job_admin.',
'job_blue-collar', 'job_entrepreneur', 'job_housemaid', 'job_management', 'job_retired', 'job_self-
employed', 'job_services', 'job_student', 'job_technician', 'job_unemployed', 'job_unknown',
'marital_divorced', 'marital_married', 'marital_single', 'marital_unknown', 'education_basic.4y',
'education_basic.6y', 'education_basic.9y', 'education_high.school', 'education_illiterate',
'education_professional.course', 'education_university.degree', 'education_unknown', 'default_no',
'default_unknown', 'default_yes', 'housing_no', 'housing_unknown', 'housing_yes', 'loan_no',
'loan_unknown', 'loan_yes', 'contact_cellular', 'contact_telephone', 'month_apr', 'month_aug',
'month_dec', 'month_jul', 'month_jun', 'month_mar', 'month_may', 'month_nov', 'month_oct',
'month_sep', 'day_of_week_fri', 'day_of_week_mon', 'day_of_week_thu', 'day_of_week_tue',
'day_of_week_wed', 'poutcome_failure', 'poutcome_nonexistent', 'poutcome_success']]

```

```
#X = data[['age', 'campaign', 'pdays', 'previous', 'no_previous_contact', 'not_working', 'job_admin.',
'job_blue-collar', 'job_entrepreneur', 'job_housemaid', 'job_management', 'job_retired', 'job_self-
employed', 'job_services', 'job_student', 'job_technician', 'job_unemployed', 'job_unknown',
'marital_divorced', 'marital_married', 'marital_single', 'marital_unknown', 'education_basic.4y',
'education_basic.6y', 'education_basic.9y', 'education_high.school', 'education_illiterate',
'education_professional.course', 'education_university.degree', 'education_unknown', 'default_no',
'default_unknown', 'default_yes', 'housing_no', 'housing_unknown', 'housing_yes', 'loan_no',
'loan_unknown', 'loan_yes', 'contact_cellular', 'contact_telephone', 'month_apr', 'month_aug',
'month_dec', 'month_jul', 'month_jun', 'month_mar', 'month_may', 'month_nov', 'month_oct',
'month_sep', 'day_of_week_fri', 'day_of_week_mon', 'day_of_week_thu', 'day_of_week_tue',
'day_of_week_wed', 'poutcome_failure', 'poutcome_nonexistent', 'poutcome_success']]
```

```
print("X: ", X)
```

```
#y = data['Loan_Status_label']
```

```
y = df['Loan_Status_label']
```

```
print("y: ", y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 42)
```

```
X_train
```

```
X_test
```

```
#Feature Scaling
```

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
sc = StandardScaler()
```

```
X_train_scale = sc.fit_transform(X_train)
```

```
X_test_scale = sc.transform(X_test)
```

```
print(X_train_scale.shape, X_train_scale)
```

```
print(X_test_scale.shape, X_test_scale)
```

```
#Using Logistic Regression
```

```
lr = LogisticRegression()
```

```
lr.fit(X_train, y_train)
```

```
print(f'Training Accuracy - {lr.score(X_train, y_train)}')
```

```
print(f'Test Accuracy - {lr.score(X_test, y_test)}')
```

```
lr = LogisticRegression()
```

```
lr.fit(X_train_scale, y_train)
```

```
print(f'Training Accuracy - {lr.score(X_train_scale, y_train)}')
```

```
print(f'Test Accuracy - {lr.score(X_test_scale, y_test)}')
```

#Using Decision Tree

```
dt = DecisionTreeClassifier(max_depth = 4)
dt.fit(X_train, y_train)
print(f'Training Accuracy - {dt.score(X_train, y_train)}')
print(f'Test Accuracy   - {dt.score(X_test, y_test)}')

dt = DecisionTreeClassifier(max_depth = 4)
dt.fit(X_train_scale, y_train)
y_pred_dtr = dt.predict(X_train_scale)
y_pred_dts = dt.predict(X_test_scale)
accuracy_score(y_train, y_pred_dtr)
accuracy_score(y_test, y_pred_dts)
print(f'Training Accuracy - {dt.score(X_train_scale, y_train)}')
print(f'Test Accuracy   - {dt.score(X_test_scale, y_test)}')
```

#Using Random Forest

```
rf = RandomForestClassifier(max_depth = 4, random_state = 42)
rf.fit(X_train, y_train)
print(f'Training Accuracy - {rf.score(X_train, y_train)}')
print(f'Test Accuracy   - {rf.score(X_test, y_test)}')

xgb = XGBClassifier(gamma = 0.5, reg_alpha = 0.6, reg_lambda = 0.3)
xgb.fit(X_train, y_train)
print(f'Training Accuracy - {xgb.score(X_train, y_train)}')
print(f'Test Accuracy   - {xgb.score(X_test, y_test)}')
```

#Using Xgboost

```
xgb = XGBClassifier(gamma = 0.5, reg_alpha = 0.6, reg_lambda = 0.3)
xgb.fit(X_train_scale, y_train)
y_pred_xgtr = xgb.predict(X_train_scale)
y_pred_xgts = xgb.predict(X_test_scale)
confusion_matrix(y_train, y_pred_xgtr)
sns.heatmap(confusion_matrix(y_train, y_pred_xgtr), annot = True, fmt = 'd')
accuracy_score(y_train, y_pred_xgtr)
confusion_matrix(y_test, y_pred_xgts)
```

```

sns.heatmap(confusion_matrix(y_test, y_pred_xgts), annot = True, fmt = 'd')

accuracy_score(y_test, y_pred_xgts)

plt.show()

#Hyperparameter Tuning with Xgboost

#Define the parameter grid

parameters = {

    'n_estimators': [100, 200],

    'learning_rate': [0.01, 0.05],

    'max_depth': [3, 4, 5],

    'gamma': [0.2, 0.3],

    'reg_lambda': [0.1, 1],

    'reg_alpha': [0.1, 1],

}

#Perform Grid Search

grid_search = GridSearchCV(estimator=xgb, param_grid=parameters, scoring='accuracy', cv=5,
verbose=3, n_jobs=-1)

grid_search.fit(X_train, y_train)

#Best parameters

print("Best hyperparameters: ", grid_search.best_params_)

#Best Estimator

print("Best Estimator: ", grid_search.best_estimator_)

xgb = XGBClassifier(gamma=0.2, learning_rate = 0.01, max_depth=44, n_estimators=200,
reg_Alpha=1, reg_lambda=0.1)

xgb.fit(X_train_scale, y_train)

y_pred_train = xgb.predict(X_train_scale)

y_pred_test = xgb.predict(X_test_scale)

confusion_matrix(y_train, y_pred_train)

accuracy_score(y_train, y_pred_train)

confusion_matrix(y_test,y_pred_test)

accuracy_score(y_test, y_pred_test)

```